

Data Link Layer: Sliding Window Protocols

Kameswari Chebrolu

Experience

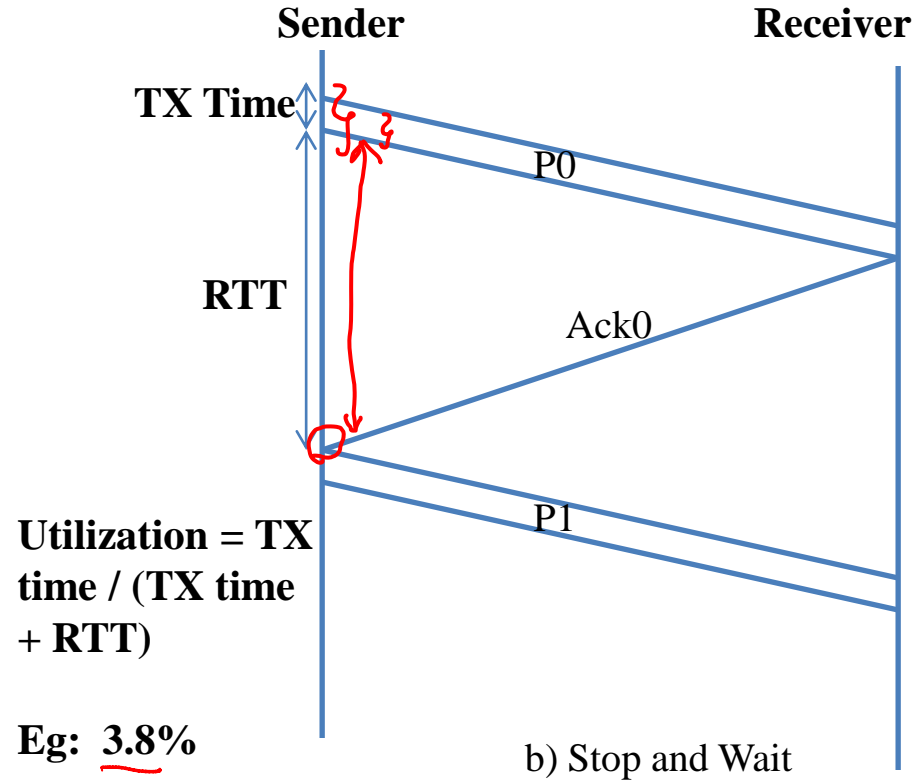
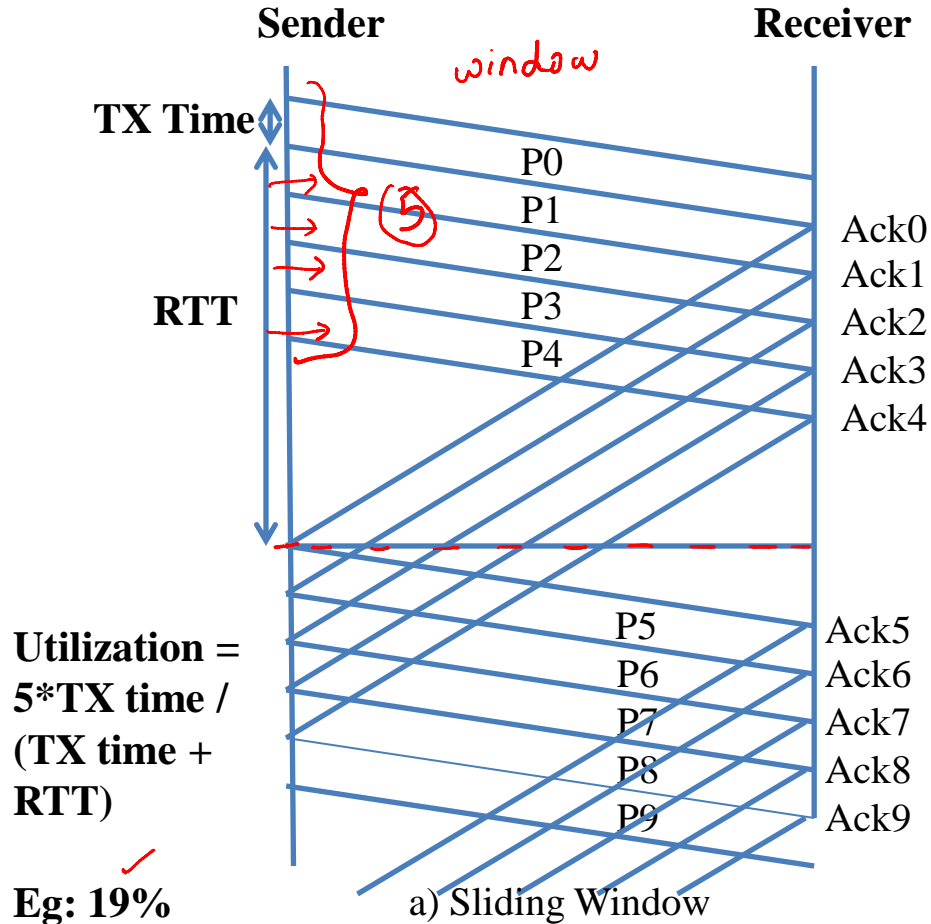
A man who carries a cat by the tail learns something he can learn in no other way.

-- Mark Twain

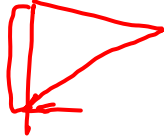
Recap

- Incrementally built the framework for a Reliable Data Transfer (RDT) protocol
 - Required Functionality
 - Stop and Wait Protocol: Works correctly but performance is poor
- Sliding Window Protocols: Can improve performance considerably

General Idea



Bandwidth-Delay Product

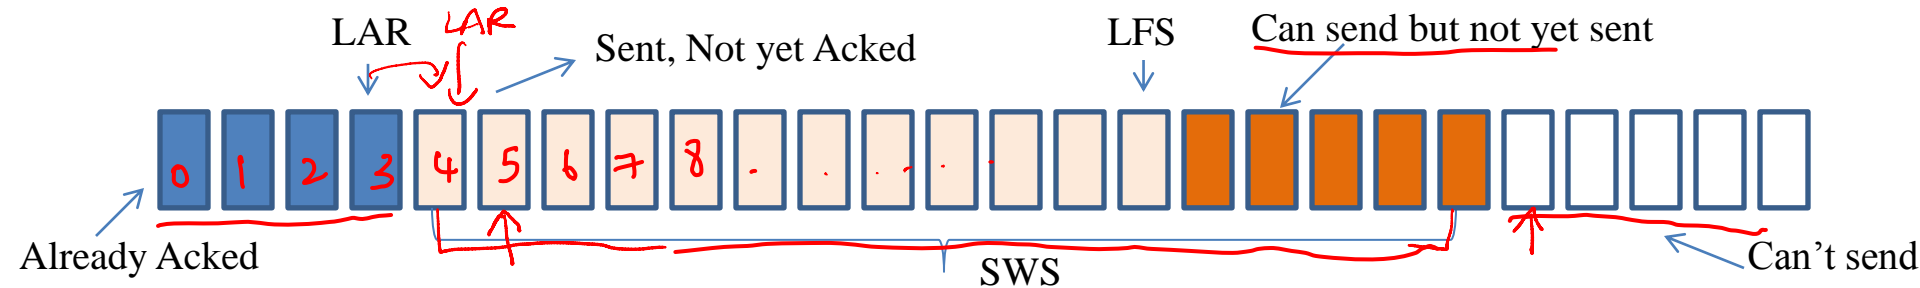
- Sender can send a maximum of W packets (window size) without waiting for an ACK
- What can the window size be?
 - To keep utilization maximum, sender can send ‘roughly’ upto Data-Rate*RTT before hearing an ack
 - Window size = Bandwidth * RTT bits
 - Divide by packet size to get the number of packets

Sliding Window

- At any given time, no more than W packets can be outstanding (their status not known at sender)
- As status of packets gets known at sender, the window of sequence number slides to encompass newer sequence numbers
 - Permits sender to send subsequent packets

Sender Side

- Assign a sequence number to each frame
- Maintain 3 variables:
 - Send Window Size (SWS): upper bound on number of outstanding frames
 - LAR: Seq. No of Last Acknowledgment Received; Advance LAR when ACK arrives
 - LFS: Sequence number of Last Frame Sent
- Maintain Invariant: LFS-LAR \leq SWS



Receiver Side

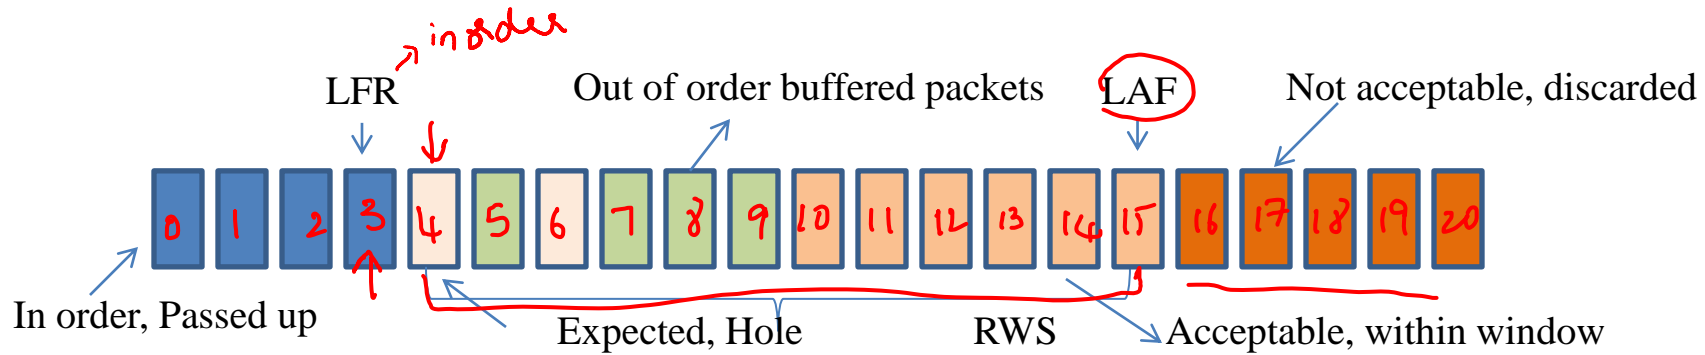
- Maintains the following three variables

- Received Window Size (RWS): upper bound on the number of out of order frames
- LAF denotes sequence number of last acceptable frame
- LFR denotes sequence number of last frame received

$RWS > SWS$?
 $RWS \leq SWS$
 $RWS = 1$?

Receiver accept only in order packets

- Set $LAF = LFR + RWS$



Consequences

$w=1$

- Need a range of sequence numbers (two won't suffice)
- Sender has to buffer more than one frame (buffer all transmitted but not yet acked)
- Receiver may also have to buffer (out of order frames) $rws \leq sws$
- Range of protocols based on sliding window
 - Go Back N, Selective Repeat, TCP
 - Actions taken on events (like frame/ack rcvd, duplicate ack, timeout etc) differ

Action at the sender

- On receiving a packet from higher layer:
 - Send the packet or buffer it or discard it?
- On receiving ACK:
 - Send packet?
 - If so, which one?
- Timeout: Retransmit packet, update timer

Action at Receiver

- On receiving a packet:
 - Should I accept it or discard?
 - Whether accepted or not, should I generate ACK?
 - If ACK generated, what should I acknowledge in it?

Receiver Side Details


- Frame “SeqNum” arrives
 - If SeqNum \leq LFR or SeqNum $>$ LAF, discard frame
 - If LFR $<$ SeqNum \leq LAF, accept
- Always good to send an ACK (improves performance by early detection of losses)
 - Exception: depends on protocol (e.g. Go-Back-N doesn't send Ack when frames are discarded)
- Types of ACK: Cumulative Ack; Selective ACK
 - Tradeoff: Simplicity vs Performance

Cumulative ACK

- SeqNumToAck
 - Largest sequence number received such that all frames with sequence number less than SeqNumToAck have also been received
 - LFR is set to this value

Example

- Frames 0,1,2,3,5,7 arrive in that order
- State: Receiver acked frame 3, assume RWS is 4
 - When 5 arrives: Buffer 5, ack 3, LFR=3, LAF=7
 - When 7 arrives, Buffer 7, ack 3, LFR=3, LAF=7
 - When 4 arrives (say due to sender retransmission), it fills the hole,
 - Accept 4 and pass up packets till 5, ack 5, LFR=5, LAF=9
 - When 6 arrives, accept/pass up, ack 7, LFR=7, LAF=11



4 frame

+RWS

Selective ACK

- For out of order packets: ACK sequence
number of whatever packet accepted
- In previous example: 0, 1, 2, 3, 5, 7
 - When 5 arrives, ack 3, sack 5
 - When 7 arrives, ack 3, sack 7 5
 - When 4 arrives, ack 5, sack 7
 - When 6 arrives, ack 7

Sender Side Details

- On receiving packet from higher layers:
 - Determine next sequence number to use (SeqNum)
 - If $\text{SeqNum} - \text{LAR} \leq \text{SWS}$, send the packet, else buffer or inform application

Sender Side Details

- On receiving ACK: action function of ACK type
 - Can get very complex (e.g. TCP variants)
 - Duplicate Acks and selective acks can be used to determine packet loss, which in turn trigger retransmission
 - Timeout is a backup mechanism → empty the pipe
 - Simple Version: Use cumulative ack to advance window
 - If SeqNum in ACK $>$ LAR, LAR = SeqNum in ACK
 - If SeqNum of Buffered packets \leq LAR+SWS, send them

Sequence Number Space

- How many sequence numbers do you need?


- Suppose RWS = 1, SWS = 3

2 seq #

W = 1


- Maxseqno = 4 i.e. 0, 1, 2, 3? 0

P0 P1 P2 P3 P4



- Suppose RWS=SWS=3

- Maxseqno=4 sufficient?



- What is the general rule?

RWS, SWS ?

sender : 0, 1, 2 next

receiver : 3, 0, 1 previous

Ack lost, Sender ret x 0, 1, 2

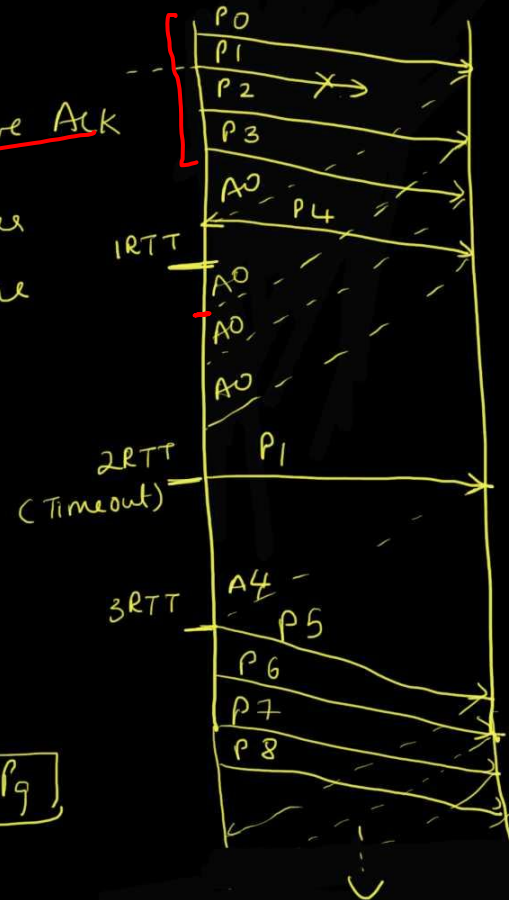
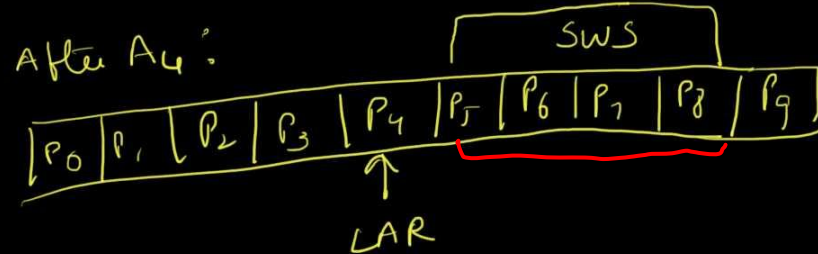
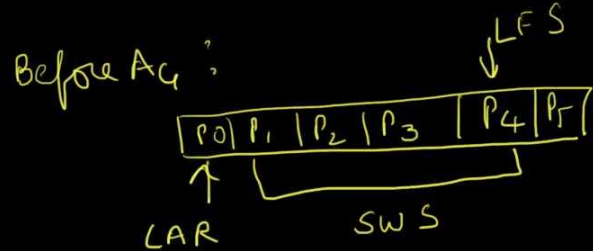
Simple Sliding Window Example

$$SWS = RWS = 4$$

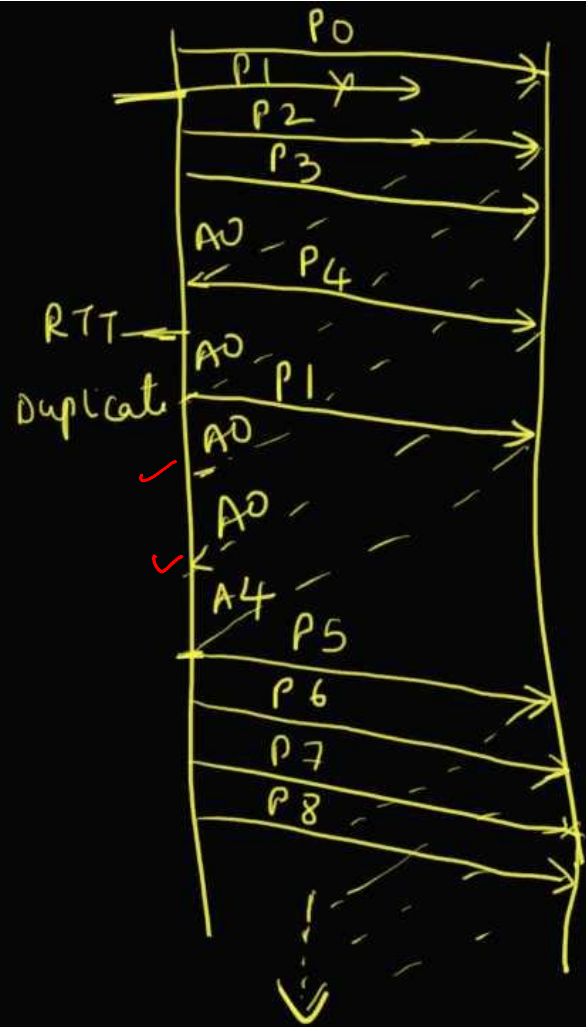
* On receipt of a frame, receiver sends cumulative Ack

* Timeout = $2 \times RTT$

* Retransmission at sender
due to timeout (ignore dup acks)



- * $SWS = RWS = 4$
- * on receipt of a frame,
receiver sends cumm. Ack
- * Timeout = 2 RTT
- * Retransmission @ Sender
due to duplicate ACK



Go Back N

state maintenance
easy

- On TimeOut, sender resends all packets that have been previously sent but unack'ed
- RWS = 1
 - Receiver accepts only in order packet and generates cumulative ack
 - Discards out of order packets, no ack generated
- Inefficient but receiver design is very simple


Go Back N

- * Receiver discards out of order packets;
NO ACK generated

- * Generator cum ACK for in order packets



- * Sender relies on Timeout for Retransmission

- * On timeout, all sent but unacked packets are sent



Selective Repeat

- Selective Repeat
- * SWS = RWS = 5
 - * Multiple losses
 - * Selective ACK



TCP

Flow Control

- Prevents the sender from overflowing the receiver buffer *small*
- Receiver informs sender how many frames it has room to receive.
 - Sender will always respect this when sending frames *SwS*

Summary

- Sliding Window protocols rely on 'window of packets' to improve performance
 - More complex (range of sequence numbers, types of acks, action at sender/receiver)
- In addition to reliability, they provide
 - In-order delivery ✓
 - Flow control ✓
- Milestone: Achieved point-to-point reliable and efficient data transfer (Building block)

