

TCP Congestion Control -- Overview

Kameswari Chebrolu

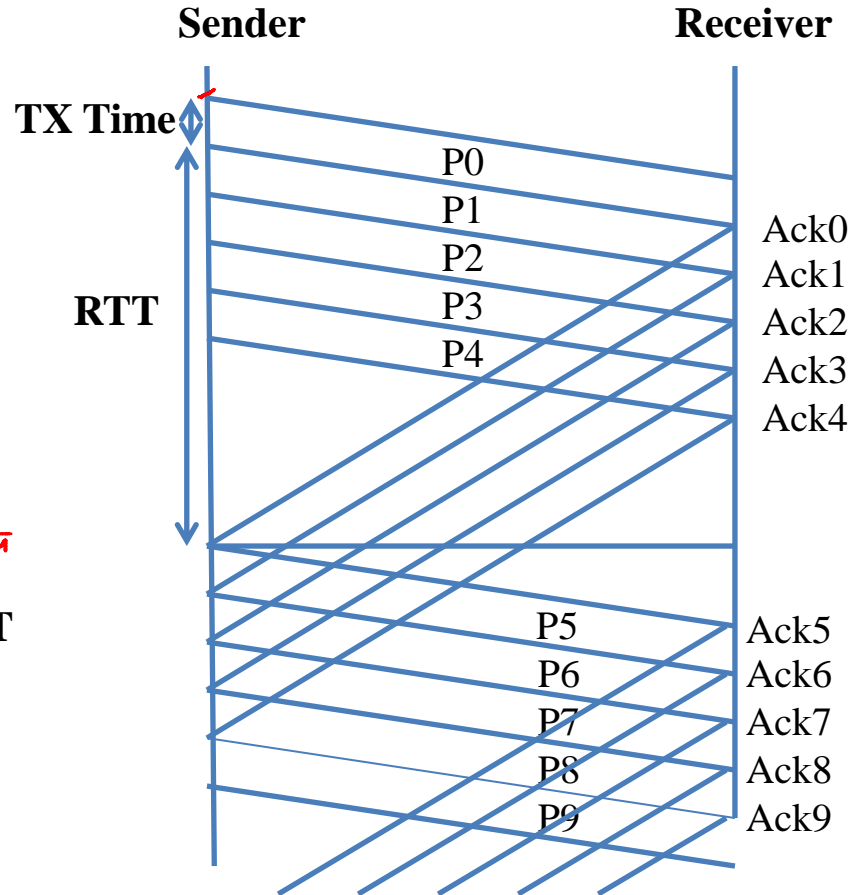
Seminal Paper: Congestion Avoidance and Control
by Van Jacobson and Michael J. Karels

Recap: TCP Services

- Multiplexing/Demultiplexing
- Reliable point-to-point data transfer
- Full-duplex
- Congestion control
- Flow control

Sliding
window
protocol

Recap: Sliding Window




Throughput $\sim (W * MSS) / RTT$

Handwritten annotations:


- A red circle around W with an arrow pointing to it and the text "wind. size" (window size).
- A red arrow pointing to MSS with the text "pkt" (packet).

Congestion Control: Problem Statement



- Many flows pass through a router; number varies with time
- Flows can be TCP or UDP ^{fairness}
- The link capacities of the routers are different
- End Result: Throughput achieved by a given flow function of many factors ^{blocks}

Congestion Control: Challenge




- Need to estimate W (of sliding window) such that each flow gets its fair share
 - Estimate small \rightarrow underutilization; Estimate large \rightarrow Congestion
- W will vary over time
- Congestion Control: Preventing sources from sending too much data too fast and thereby ‘congest’ the network

Sliding Window Protocol

- Roughly, idea translates to the following:
- View network as a pipe
- Determine the capacity of the pipe (Bandwidth-delay product)
- Fill the pipe with data
- As you remove one packet from the pipe, add another
 - ACKs help clock out data (Self Clocking)

Self-Clocking



W → window size in bits

① → $W = B \times RTT$

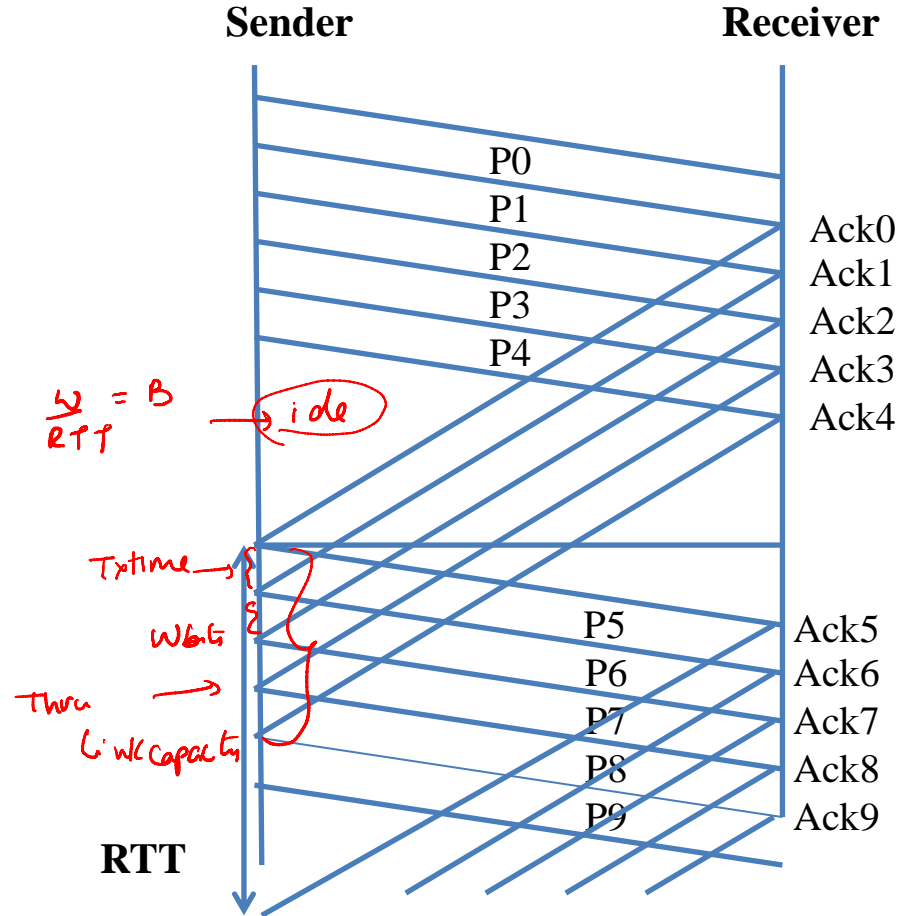
② Max Thr A & B - ? = B


③ Rate @ which Ack's are coming?
→ Rate @ which Pkts
↳ Link Capacity

$W < B \times RTT$

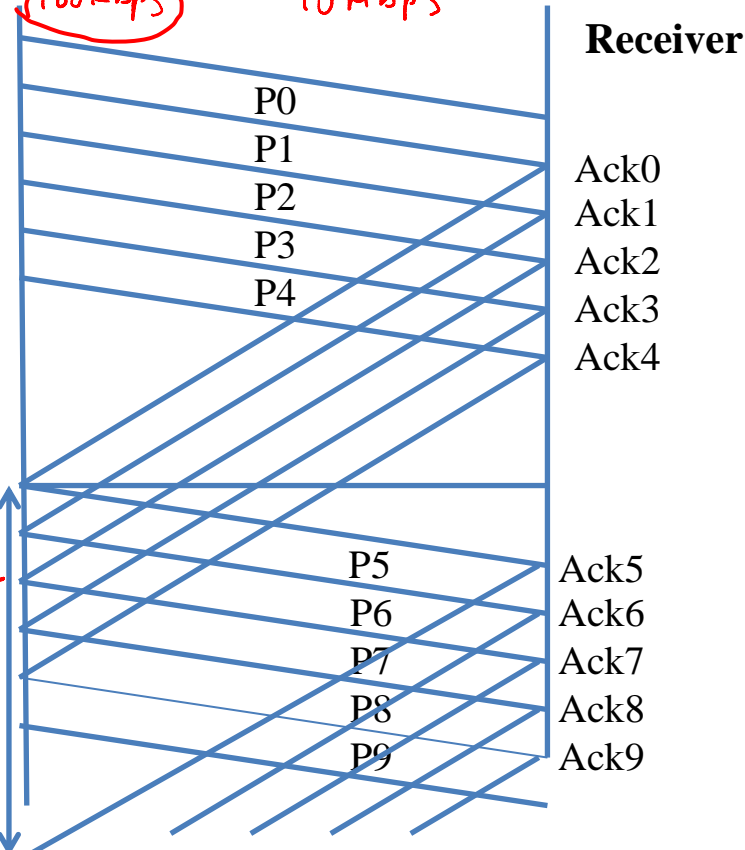
→ Link capacity


Self-clocking



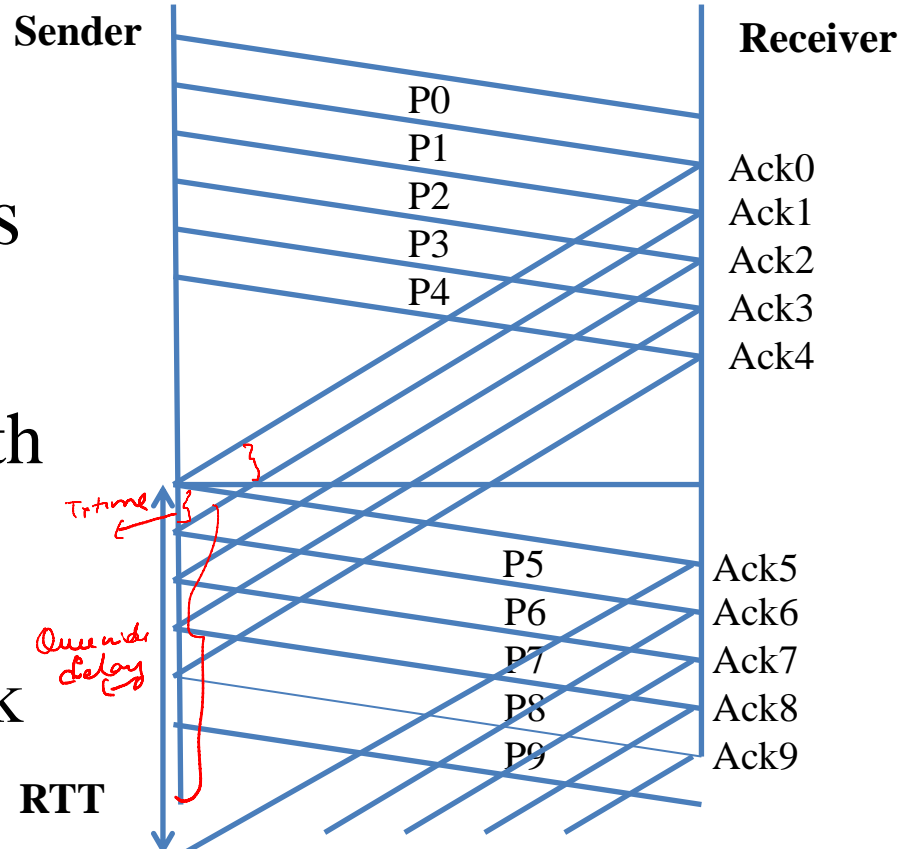
$$RTT \times 100 \text{ Kbps}$$


- ② Max Thr. = ? 10 Mbps 100Mbps **Send**
- ③ rate @ acis @ sender ?
 ↳ rate @ phts on this bottleneck link






- $W > \text{BDP}$ (in bits)
 $\hookrightarrow \text{RTT}_{\text{no load}}$
- Sliding Window says
 - $\text{Thr} = \frac{W}{\text{RTT}}$ $\hookrightarrow \text{RTT}_{\text{load}}$
Bottleneck bandwidth
- Self-clocking says
 - Max thr is bottleneck bandwidth



Window Size



Same Idea

- View network as a pipe
- Estimate Bandwidth-delay product (capacity) dynamically
 - Uses the variable Congestion Window (CW) to track it
- Self Clocking: Captures bottleneck bandwidth
 - Use ACKs to clock out data
 - Not perfect (With competing traffic ack spacing will not be preserved)

3 Steps

- Getting to Equilibrium[→]
- Conservation at equilibrium_→
 - Don't put new packet unless old one is removed
- Adapting to Path Dynamics

Summary

- Congestion Control is a complex problem
- Need to implement it in the context of the sliding window protocol
 - Self clocking[✓] is a useful feature (we will rely on this to capture bottleneck bandwidth)
 - Need to determine and adapt W (window size) such that you don't underutilize bandwidth or congest the network
- Ahead: Actual details