


TCP Congestion Control Mechanisms

Kameswari Chebrolu

Seminal Paper: Congestion Avoidance and Control
by Van Jacobson and Michael J. Karels

Congestion Control: Challenge



- Need to estimate W (of sliding window) such that each flow gets its fair share
 - Estimate small \rightarrow underutilization; Estimate large \rightarrow Congestion
- W will vary over time
- Congestion Control: Preventing sources from sending too much data too fast and thereby ‘congest’ the network

Idea

- View network as a pipe
- Estimate Bandwidth-delay product (capacity) dynamically
 - Uses the variable Congestion Window (CW) to track it
- Use self clocking to pump packets into the network


Approach

- Getting to Equilibrium *→ filling up the pipe*
- Conservation at equilibrium
- Adapting to Path

Getting to Equilibrium

BDP - (1) (10, 100, 200)
↑


- What value of CW to choose initially?
 - Too large: pushes network into congestion
 - Just right: bursty transmissions can lead to losses



Slow Start


- Add a variable cwnd (congestion window)
 - Captures the number of outstanding data in the network
- At start, set cwnd=1
- On each ack for new data, increase cwnd by 1

- Takes $\log_2 W$ round trip times to get to W
- Sends data at most twice the bottleneck link on the path
- Bound to overestimate capacity at some time in future



Conservation at Equilibrium

- Don't put a packet unless a packet is removed
 - Particularly important when the network is congested
 - Can potentially happen on timeouts → proper RTT estimation crucial
 - Delayed packets should not be interpreted as lost



RTT Estimation: Original Algorithm

- Measure SampleRTT for sequence/ack combo
- $\text{EstimatedRTT} = a * \text{EstimatedRTT} + (1 - a) * \text{SampleRTT}$
 - Small a heavily influenced by temporary fluctuations
 - Large a not quick to adapt to real changes
 - a is between 0.8-0.9
- $\text{Timeout} = 2 * \text{EstimatedRTT}$
 - variability* ←


Jacobson/Karels Algorithm

- Algorithm takes into account variance of RTTs
 - If variance is small, EstimatedRTT can be trusted
 - If variance is large, timeout should not depend heavily on EstimatedRTT

- $\text{Difference} = \text{SampleRTT} - \text{EstimatedRTT}$
- $\text{EstimatedRTT} = \text{EstimatedRTT} + (d * \text{Difference})$
- $\text{Deviation} = \text{Deviation} + d (|\text{Difference}| - \text{Deviation})$, where $d \sim 0.125$
- $\text{Timeout} = u * \text{EstimatedRTT} + q * \text{Deviation}$,
where $u = 1$ and $q = 4$
- Exponential Timeout backoff: controls spacing between retransmits



Exponential Damping

- From control theory: An unstable system can be stabilized by adding exponential damping
- “A network subject to random load shocks and prone to congestive collapse can be stabilized by adding exponential damping to its primary excitation (Traffic sources)”


Adapting to Path

- Estimating process can over or underestimate W ; need to correct this
- Available bandwidth also changes over time; need to adapt to this
- Need a feedback mechanism from the network that the estimate is wrong

Overestimation

- Overestimation leads to congestion
- Feedback: If losses are due to congestion and ^{TO} timers are working correctly → Timeout indicates congestion
- How to change the congestion window? $(w) - c$
 $(TO) \rightarrow$
 w/c
 - Additive decrease or multiplicative decrease?
 - Multiplicative decrease yields better stability
 - $W_i = dW_{i-1}$ ($d < 1$, typically 0.5)
 $w = 10$
 $w = 5$

Underestimation

$$w = 5 \xrightarrow{+7} 15$$

- Underestimation leads to lower utilization
- Additive increase or multiplicative increase?
 - Exponential increase leads to instability; overestimation is inevitable
 - Additive increase
 - $W_i = W_{i-1} + u$ ($u \ll W_{\max}$; typical u is 1)
 - Increase window by 1 segment every RTT

$$w \neq c$$


$$c \times w$$

Congestion Avoidance

- Additive Increase, Multiplicative Decrease
- On detecting congestion, set cwnd to half the window size (multiplicative decrease)

- *slow down &* On each ack of new data, increase cwnd by $1/\text{cwnd}$ (additive increase)

RTT \rightarrow cwnd by 1 segment \checkmark



Summary

- Congestion control is a difficult task
 - Prevent underutilization; ensure no congestion; ensure fairness
- TCP relies on a variety of techniques to achieve this
 - Slow start, RTT estimation, Congestion avoidance (AIMD)
- Ahead: Putting it all together in TCP versions